

Attacking Kubernetes Clusters

- An attack on Kubernetes generally starts from the perspective of a compromised pod.
- The threat actor reaches this point via a scenario similar to these:
 - Actor compromised the application running in one container in the pod.
 - Actor phished/compromised a person who had access to the pod.
 - Actor was authorized and wants to escalate their privileges.
- As a defender, once you can handle the compromised pod scenario, it's time to gain the ability to handle a compromised node.
 - Nodes are compromised either directly, through phishing/social engineering attacks, or through container breakouts.

Attacks from within a Compromised Pod

An attacker in a pod may, among other things:

- Use the access provided by the pod to access other services`
- Attack other containers in their pod
- Make requests to the API server or a Kubelet to:
 - Run commands (possibly interactively) in a different pod
 - Start a new pod with privilege and node filesystem/resource access
 - Gather secrets that Kubernetes provides to pods
- Connect to the Kubernetes dashboard to perform actions
- Interact with the etcd server to change the cluster state
- Interact with the cloud service provider using a cluster account.

Microsoft's Threat Matrix for Kubernetes

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking	Denial of service
Application vulnerability	Application exploit (RCE)		Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files	
Exposed Dashboard	SSH server running inside container					Instance Metadata API	Writable volume mounts on the host	
							Access Kubernetes dashboard	
							Access tiller endpoint	

Defense: Overarching Note

You must upgrade your Kubernetes cluster.

Kubernetes development is active and moves very quickly.

The Kubernetes project supports only the last year's worth of releases. If a cluster is more than 12 months old, it may very well no longer have security patches available.

Before the third quarter of 2020, Kubernetes only supported 9 months of releases.

Additionally, the Kubernetes security defaults and capabilities continue to improve.

Defense: RBAC and Authorization (Authz)

- Role-based Access Control (RBAC)
- Removing default service account permissions

Role-Based Access Control

- You can place restrictions on the API server via RBAC.
- RBAC defines what PRINCIPALS can perform what ACTIONS.
- Principals are users or service accounts.
 - Example: [jay in group system:authenticated]
- Actions are VERBS combined with OBJECT types:
 - Example: [create namespace]
 - Example: [in a specific namespace, create apps deployment]

Role-Based Access Control: Roles

- You provide the ability to do these things by creating:
 - Role – specifying a list of actions
 - Role Binding – allowing a principal to use a role (list of actions).
- Roles have a many-to-many relationship with principals.
- Roles and Role Bindings are scoped to a namespace.
 - To scope globally, use Cluster Roles and Cluster Role Bindings.

Create a Role and RoleBinding

```
kind: Role
apiVersion: ...
metadata:
  name: ing-pod-reader
  namespace: inguardians-ns
rules:
- verbs: ["get","list"]
  apiGroups: [""]
  resources: ["pods"]
```

```
kind: RoleBinding
apiVersion: ...
metadata:
  name: frontend-pod-reader
  namespace: inguardians-ns
roleRef:
  kind: Role
  apiGroup: ...
  name: ing-pod-reader
Subjects:
- kind: ServiceAccount
  apiGroup: ...
  name: frontend
```


Creating Custom Roles Automatically

Jordan Liggitt wrote a tool called Audit2RBAC, similar to Audit2Allow for SELinux.

<https://github.com/liggitt/audit2rbac/>

Watch this in action via this video:

<https://www.youtube.com/watch?v=n2cD20moYe8&feature=youtu.be>

Default Service Account Permissions

Once you have custom service accounts defined and working, remove permissions on the default service accounts.

Reference:

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

Exercise: Kubernetes Own the Nodes

We're going to do our first Kubernetes exercise now.

Please:

Open the Firefox browser on the class machine to:

<http://localhost:10000/exercises/kubernetes-own-the-nodes>

You may need to replace your daemonset file. Please find a replacement in the repo.

Replacement daemonset

containers:

- name: attack-root

image: k8s.gcr.io/redis:e2e

resources:

limits:

memory: 100Mi

cpu: 1m

requests:

cpu: 1m

memory: 100Mi